

An Application Management Tool in Go

Marc Sauter, René Zbinden PostFinance



A Brief History of Application Management Tools @PostFinance

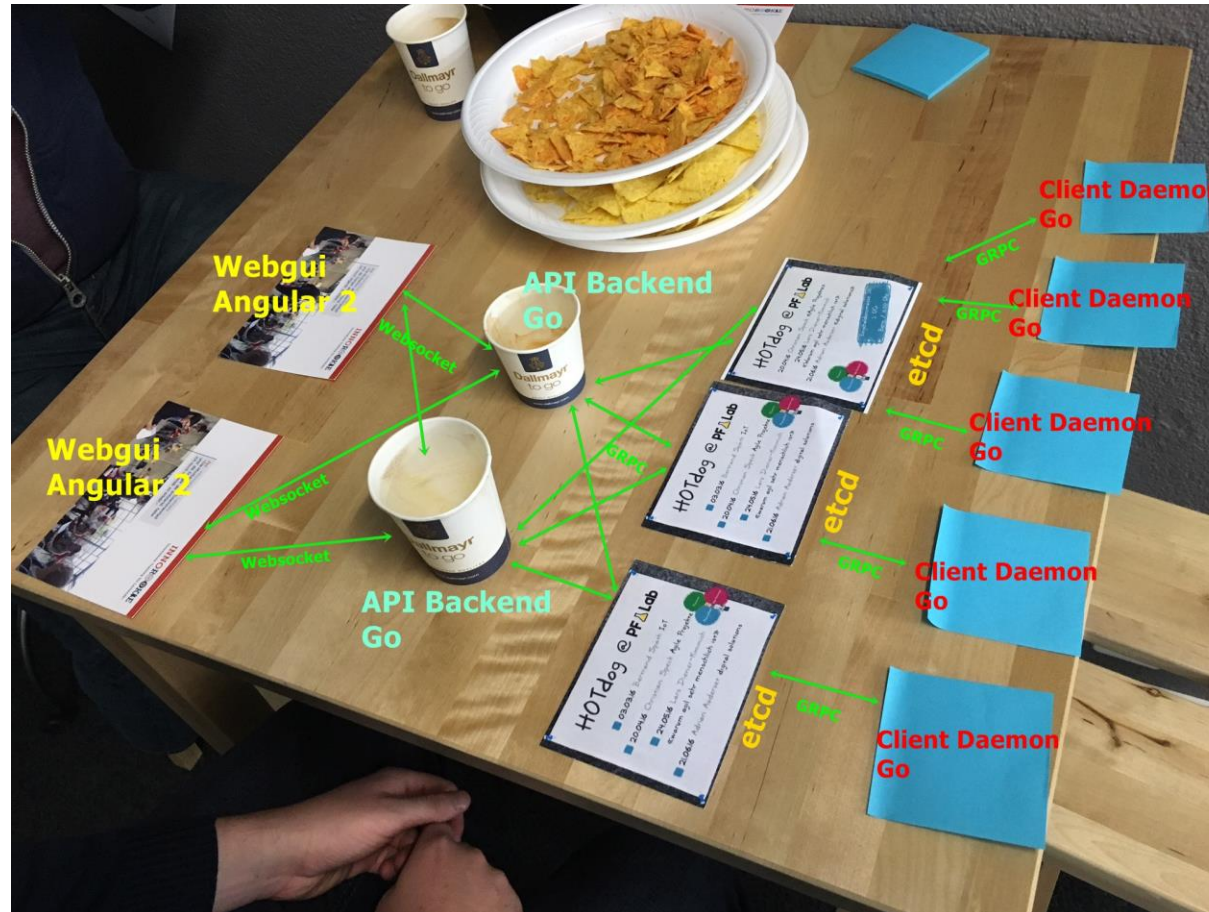

- C -> Perl (TUI)
- TCL (TUI)
- Java (Agent, Backend, GUI)
 - Performance problems -> (very expensive) refactoring -> still problems -> project terminated
- Revival of Perl due to the lack of alternatives
- Mule ... since 2016 and still no need to replace it!
 - the first tool where the stakeholders were able to configure it





Go@PostFinance

- 2014 (maybe even 2012) First command line tools
- 2015 Certificate Enrollment System
 - Agent for Linux, Solaris and Windows
- 2016 Mule
 - Backend on Kubernetes
 - Agent for Linux, Solaris and Windows
- Today Go is the first choice for infrastructure tools and IaC.
 - PostFinance uses mainly Java and Oracle for application development.

First Design - 21.04.2016



Architecture





Components


- Backend
 - GUI-Server (REST-API)
 - Go
- Agent
 - Go
- GUI
 - Single Page Application with ReactJS
- Store
 - etcd
- Configuration
 - YAML files in a Git repository ... YAML was a well known format for the stakeholders



Principles

- Configuration in YAML
 - Easy to read/edit
 - Self Service
- Applications and Services get discovered and registered automatically
- The runtime state can be recovered from the configuration
 - etcd backup is not necessary
- At most one writer per etcd key to prevent concurrency issues
- The stakeholders are part of the project

GUI – developed by the stakeholders

[illegible]



Launch

- It worked much better than we expected.
- It scaled to > 2k agents.
- But we still had to make a few pprof flamegraphs to fix some performance issues.
- Many performance issues in the beginning were related to etcd. Most issues could be fixed with code optimization, some were fixed with etcd upgrades.
- Added the etcd-proxy to reduce the number of connections to the etcd itself.

Go was the right choice!



Problems

- Code Quality
 - After 5 years of extending and bugfixing, the code was not in good shape anymore.
 - Mule is a kind of hobby and not our main job.
- ETCD
 - Etcd was the core of Mule. Every server and agent instance was connected to etcd and had multiple watchers running.
 - Etcd was used as a runtime store and (misused as) a message queue.
 - > 2k active clients
 - > 4k watchers
 - very sensitive to network and/or disk I/O latency



Refactoring

1. Refactor code without changing the architecture
2. Review Architecture
3. Refactor code with new architecture




Refactoring – Keep Architecture

- Configure `golangci-lint`
- Rewrite discovery service
- Improve Observability
- Replace
 - `pkg/errors` with `errors` from `stdlib`
 - custom logger implementation with `uber-go/zap`


Changes 534

Showing 20 changed files ▾ with 1527 additions and 462 deletions

NATS - Architecture



NATS




- Released April 2022

Changes 121

Pipelines 1

Showing 20 changed files ▾ with 615 additions and 1104 deletions

NATS





Several «Small» Bugs

feat: better approval validation ...

René Zbinden authored 1 year ago

fix: only debug log when no output found ...

René Zbinden authored 1 year ago

fix: set lifecycle data for stopped or failed nodes ...

René Zbinden authored 1 year ago

fix(agent): set LimitNOFILE=16384 in systemd

René Zbinden authored 1 year ago

fix: only log action for current node ...

René Zbinden authored 1 year ago

```
{"level":"INFO","ts":"2023-05-02T13:29:18.349+0200","caller":"repo/repo.go:177","msg":"put object","key":"ins
initiation":"systemd","name":"chronyd.service","current-status":"running","target-status":"running","override-st
panic: assignment to entry in nil map
goroutine 259 [running]:
gitlab.pnet.ch/mule/mule/internal/metrics.ProfilesCollector.collectEmptyProfiles({0xc000012390, 0xc000480f80,
/appl/sdsa/gitlab-runner/builds/eqWDYyC6/0/mule/mule/internal/metrics/profiles.go:279 +0x4ed
gitlab.pnet.ch/mule/mule/internal/metrics.ProfilesCollector.Start({0xc000012390, 0xc000480f80, 0xc0001246f0,
/appl/sdsa/gitlab-runner/builds/eqWDYyC6/0/mule/mule/internal/metrics/profiles.go:96 +0x15e
created by gitlab.pnet.ch/mule/mule/internal/http/api.(*API).startMetricsCollector
/appl/sdsa/gitlab-runner/builds/eqWDYyC6/0/mule/mule/internal/http/api/api.go:491 +0x1ca
```




One Big Bug

- June 2022 (after several weeks in production)
- Incident: The GUI is unresponsive
- Server stopped handling Jetstream messages
- No Errors in Logs
- OS Patch of NATS Servers





Select isn't Broken

Programmers will sometimes think the problems they encounter are the fault of bugs in well-tested system software, rather than their own code.

The Pragmatic Programmer, by Hunt and Thomas

Pending Messages




Bug



```
1 go func() {
2     for {
3         ctx, cancel := context.WithTimeout(pctx, maxWait)
4
5         msgs, err := s.Fetch(1, nats.Context(ctx))
6
7         cancel()
8
9         if err != nil {
10             break
11         }
12
13         q.handleMsg(msgs[0])
14     }
15 }()
```

Bug Fix



```
1 go func() {  
2     for {  
3         ctx, cancel := context.WithTimeout(pctx, maxWait)  
4  
5         msgs, err := s.Fetch(1, nats.Context(ctx))  
6  
7         cancel()  
8  
9         if err != nil {  
10             l.Errorw("failed to get next message", "err", err)  
11  
12             continue  
13         }  
14  
15         q.handleMsg(msgs[0])  
16     }  
17 }()
```



Refactorings in the Future

- Replace `uber-go/zap` logger with `log/slog` from `stdlib`
- Get rid of `google.golang.org/grpc` dependency
 - Replace with `connectrpc/connect-go`
 - Built on `net/http`
 - Automatic protocol detection
 - Semantic versioning
 - `grpc-web` support
- Tracing?



Conclusion

- Go was the right choice.
 - Static type system (refactoring)
 - Easy cross compilation (Windows and even Solaris)
 - Static binaries for easy deployment
- NATS was the right choice.
 - Synthetic load test with ~10k clients
- It's hard to predict the load of a new distributed system, with unknown load requirements.



Why NATS?

- Simplicity (easy to deploy and monitor)
- High Performance and low latency
- Scalability and resilience (scale horizontally, built-in fault tolerance)
- Many communication patterns (request-reply)
- Testability

```
func TestWithNats(t *testing.T) {  
    start := time.Now()  
    _ = natstest.RunRandClientPortServer()  
    fmt.Println(time.Since(start))  
}
```




Minor Issues with NATS

```
> 30.09.23 23:59:59,000 Sep 30 23:59:59 [1] nats-server[1414]: :6222 - rid:30635085 - TLS route handshake error: x509: cannot validate certificate for because it doesn't contain any IP SANs
host = source = /var/log/messages sourcetype = syslog

> 30.09.23 23:59:59,000 Sep 30 23:59:59 [1] nats-server[1414]: :6222 - rid:30635084 - Router connection closed: TLS Handshake Failure
host = source = /var/log/messages sourcetype = syslog
```

Server Overview

Host	Version	JS	Conns	Subs	Routes	GWs	Mem	CPU %	Cores	Uptime	RTT
	2.9.17	yes	1,343	8,195	2	0	323 MiB	3	16	70d23h29m12s	1ms
	2.9.17	yes	1,183	8,037	2	0	283 MiB	3	16	70d22h28m14s	3ms
	2.9.17	yes	819	7,678	2	0	213 MiB	2	16	52d4h51m29s	3ms
3		3	3,345	23,910			819 MiB		3		



Fixed with 2.10.x



Refactoring - NATS

- Replaces the communication between backend and agent
- Etcd is the only holder of the runtime state for the GUI
- JetStream for state changes
- NATS for the action queue and configuration requests

Statistics





NATS - Architecture

