



Roman Scharkov

Backend Engineer

at  **tutti.ch**



swiss
marketplace
group

Real Estate



Automotive



General Marketplaces



Finance & Insurance



JSCAN^{v3}

JSON validating • JSON parsing • JSON decoding

HISTORY



(5y ago) Aug 5th 2019: github.com/romshark/llparser



(5y ago) some day 2019: started to love JSON



(3y ago) Jan 8th 2022: **jscan first commit**



(1y ago) Mar 28th 2023: **jscan v1.2 (RFC8259**

Compliance)



(9m ago) Jun 12th 2023: **jscan v2 released (perf &**


feat)



(5m ago) Oct 3rd 2023: Zürich Gophers Meetup

I'M FINISHED




A close-up shot of a woman with short, light-colored hair, looking down and slightly to the left with a somber expression. The lighting is dim and warm, creating a melancholic atmosphere. The background is dark and out of focus.

I'm not worth it.

HAND-ROLLING JSON PARSERS IS A LIABILITY



**I ONLY HAVE
UNTAPPED POTENTIAL**



REQUIREMENTS



No code generation.



Easy to use **Unmarshal** and **Decoder**.



Backward-compatibility with **encoding/json**, ideally a drop-in replacement.



Extensibility to the feature set of **json v2**
(github.com/go-json-experiment/json)



Best in class **performance** through memory trade-off

**JUST
DO IT!**



DO IT. JUST DO IT.



Nothing is impossible.



Don't let your dreams be dreams.



```
1  `{  
2      "name": "Test name",  
3      "number": 100553,  
4      "tags": ["sports", "portable"]  
5  }`
```

```
1  type Struct3 struct {  
2      Name    string    `json:"name"`  
3      Number  int         `json:"number"`  
4      Tags    []string   `json:"tags"`  
5  }
```

~**5x** faster than **encoding/json**
28% faster than fastest code generator **easyjson**
23% faster than fastest unmarshaler **goccy**
least memory allocations of all.


1	BenchmarkDecodeStruct3			
2	hand/jscan	348 ns/op	48 B/op	3 allocs/op
3	unmr/jscan	403 ns/op	32 B/op	3 allocs/op
4	unmr/goccy	526 ns/op	192 B/op	5 allocs/op
5	genr/easyjson	564 ns/op	96 B/op	4 allocs/op
6	unmr/jsoniter	640 ns/op	128 B/op	6 allocs/op
7	hand/gjson	949 ns/op	352 B/op	4 allocs/op
8	unmr/jsonv2	962 ns/op	96 B/op	3 allocs/op
9	unmr/segmentio	967 ns/op	240 B/op	5 allocs/op
10	genr/ffjson	1175 ns/op	456 B/op	11 allocs/op
11	hand/fastjson	1522 ns/op	1600 B/op	14 allocs/op
12	unmr/encoding_json	2169 ns/op	352 B/op	11 allocs/op

~**10x** faster than **encoding/json**
44% faster than fastest code generator **easyjson**
55% faster than fastest unmarshaler **goccy**.

1	BenchmarkDecodeArrayString			
2	hand/jscan	9107 ns/op	5720 B/op	25 allocs/op
3	unmr/jscan	16261 ns/op	22384 B/op	219 allocs/op
4	genr/easyjson	29354 ns/op	25944 B/op	201 allocs/op
5	unmr/goccy	36849 ns/op	24810 B/op	197 allocs/op
6	genr/ffjson	42601 ns/op	26821 B/op	210 allocs/op
7	hand/fastjson	50332 ns/op	91480 B/op	214 allocs/op
8	unmr/jsoniter	56687 ns/op	32259 B/op	321 allocs/op
9	unmr/jsonv2	74227 ns/op	21099 B/op	146 allocs/op
10	hand/gjson	76219 ns/op	75384 B/op	59 allocs/op
11	unmr/encoding_json	177211 ns/op	28888 B/op	230 allocs/op
12	unmr/segmentio	203192 ns/op	41472 B/op	337 allocs/op

not always fair to compare
like with **struct { X,Y,Z float64 }**


1	BenchmarkDecodeStructVector3D on Apple M1				
2	genr/easyjson	107.7	ns/op	0 B/op	0 allocs/op
3	hand/jscan	132.4	ns/op	0 B/op	0 allocs/op
4	unmr/goccy	161.9	ns/op	72 B/op	2 allocs/op
5	unmr/segmentio	211.5	ns/op	24 B/op	1 allocs/op
6	hand/gjson	225.2	ns/op	48 B/op	1 allocs/op
7	unmr/jscan	250.2	ns/op	24 B/op	1 allocs/op
8	unmr/jsoniter	309.8	ns/op	56 B/op	3 allocs/op
9	unmr/jsonv2	311.6	ns/op	24 B/op	1 allocs/op
10	hand/fastjson	388.5	ns/op	776 B/op	7 allocs/op
11	genr/ffjson	413.8	ns/op	352 B/op	6 allocs/op
12	unmr/encoding_json	634.9	ns/op	240 B/op	5 allocs/op
13					




```
1 j := `[1,2,3,4,5]`  
2 var v [3]int  
3 if err := json.Unmarshal([]byte(j), &v); err != nil {  
4     panic(err)  
5 }  
6 fmt.Println(v)
```

- A) [1,2,3]
- B) [3,4,5]
- C) Panic

```
1 j := `[1,2,3,4,5]`  
2 var v [3]int  
3 if err := json.Unmarshal([]byte(j), &v); err != nil {  
4     panic(err)  
5 }  
6 fmt.Println(v)
```


- A) [1,2,3] 
- B) [3,4,5]
- C) Panic

For arrays,
encoding/json **takes**
first N values and
ignores excess values.

```
1  j := `[1,null,3]`  
2  v := [3]int{7, 8, 9} // Pre-initialized!  
3  if err := json.Unmarshal([]byte(j), &v); err != nil {  
4      panic(err)  
5  }  
6  fmt.Println(v)
```

- A) [0,0,0]
- B) [1,0,3]
- C) [1,8,3]
- D) Panic

```
1 j := `[1,null,3]`
2 v := [3]int{7, 8, 9} // Pre-initialized!
3 if err := json.Unmarshal([]byte(j), &v); err != nil {
4     panic(err)
5 }
6 fmt.Println(v)
```


- A) [0,0,0]
- B) [1,0,3]
- C) [1,8,3] 
- D) Panic

null initializes **new** values to **zero value**, but **existing** values remain **untouched**.

```
1 j := `[1,2,3,"42"]`  
2 var v [3]int  
3 if err := json.Unmarshal([]byte(j), &v); err != nil {  
4     panic(err)  
5 }  
6 fmt.Println(v)
```

- A) [1,2,3]
- B) [0,0,0]
- C) [42,2,3]
- D) Panic

```
1 j := `[1,2,3,"42"]`  
2 var v [3]int  
3 if err := json.Unmarshal([]byte(j), &v); err != nil {  
4     panic(err)  
5 }  
6 fmt.Println(v)
```

- A) [1,2,3] 
- B) [0,0,0]
- C) [42,2,3]
- D) Panic

encoding/json **doesn't**
care about the **excess**
value types.

```
1 j := `{ "foo":1, "FOO":2 }`  
2 var v struct { Foo int `json:"foo"` }  
3 if err := json.Unmarshal([]byte(j), &v); err != nil {  
4     panic(err)  
5 }  
6 fmt.Println(v)
```

- A) {1}
- B) {2}
- C) Panic


```
1 j := `{ "foo":1, "FOO":2 }`  
2 var v struct { Foo int `json:"foo"` }  
3 if err := json.Unmarshal([]byte(j), &v); err != nil {  
4     panic(err)  
5 }  
6 fmt.Println(v)
```

- A) {1}
- B) {2} ✓**
- C) Panic

encoding/json **prefers exact**
matches over **case-insensitive**
matches and **overwrites** on
collision.

```
1  type User struct {
2      ID    int    `json:"id"`
3      Name string `json:"name"`
4  }
5  j := `{"id": 500}`
6  v := User{ID: 404, Name: "Bob"}
7  if err := json.Unmarshal([]byte(j), &v); err != nil {
8      panic(err)
9  }
10 fmt.Println(v)
```

- A) {500 Bob}
- B) {500 }
- C) Panic

```
1  type User struct {
2      ID    int    `json:"id"`
3      Name string `json:"name"`
4  }
5  j := `{"id": 500}`
6  v := User{ID: 404, Name: "Bob"}
7  if err := json.Unmarshal([]byte(j), &v); err != nil {
8      panic(err)
9  }
10 fmt.Println(v)
```

A) {500 Bob} 

B) {500 }

C) Panic

encoding/json
**doesn't overwrite
structs**, it mutates
the struct instead.

```
1 //      I      D
2 j := `{"\u0049\u0044":42}`
3 var v struct{ ID int }
4 if err := json.Unmarshal([]byte(j), &v); err != nil {
5     panic(err)
6 }
7 fmt.Println(v)
```

- A) {42}
- B) {}
- C) Panic

```
1 //      I      D
2 j := `{"\u0049\u0044":42}`
3 var v struct{ ID int }
4 if err := json.Unmarshal([]byte(j), &v); err != nil {
5     panic(err)
6 }
7 fmt.Println(v)
```

A) {42} 

B) {}

C) Panic

encoding/json
supports escaped
field names

```
1  type Stats struct {
2      ID    string `json:"id"`
3      Score int    `json:"score"`
4  }
5  j := `{"Bob": {"score": 4}, "Alice": {"score": 5}}`
6  v := map[string]Stats{"Bob": {ID: "b"}}
7  if err := json.Unmarshal([]byte(j), &v); err != nil {
8      panic(err)
9  }
10 fmt.Println(v)
```

- A) map[Alice:{ 5} Bob:{"b" 4}]
- B) Panic

```
1  type Stats struct {
2      ID    string `json:"id"`
3      Score int    `json:"score"`
4  }
5  j := `{"Bob": {"score": 4}, "Alice": {"score": 5}}`
6  v := map[string]Stats{"Bob": {ID: "b"}}
7  if err := json.Unmarshal([]byte(j), &v); err != nil {
8      panic(err)
9  }
10 fmt.Println(v)
```

A) map[Alice:{ 5} Bob:{"b" 4}]

B) Panic

C) map[Alice:{ 5} Bob: { 4}] 

encoding/json
always
overwrites
map keys

Not always fair to compare performance because **easyjson** and many other are **not backward compatible** with encoding/json

- no case-insensitive matching
- no unescaping
- no support for non-struct types
- etc.

Best to compare with **goccy** and **jsoniter**.

Reusing memory is key

```
1 // T is the type we want unmarshal to.
2 type T struct { Foo string; Bar int }
3
4 tokenizer := jscan.NewTokenizer[string](
5     16, // Stack size, maximum document depth.
6     1024, // Preallocated tokens buffer size.
7 )
8
9 // Preallocate decoder parsing string into `T`.
10 decoderT, err := jscandec.NewDecoder[string, T](
11     tokenizer, jscandec.DefaultInitOptions,
12 )
13 if err != nil {
14     panic(err)
15 }
16
17 // Now we can reuse decoderT many times.
18 var v T
19 if _, err := decoderT.Decode(in, &v, jscandec.DefaultOptions); err != nil {
20     panic(err)
21 }
```

Init-Time
(ones per type)

create decoding table
using reflect

Runtime
(per input)

Tokenize input



init decoding table



decode to variable

```
type T struct { IDs []int; Name string; Map map[string]float64 }
```

type ExpectType	offset,size uintptr	dest unsafe.Pointer	parent uint32	fields []fieldFrame
Struct	0,48	nil	4294967295	[{1 "IDs"}, {3 "Name"}, {4 "Map"}]
Slice	0,24	nil	0	nil
Int	0,8	nil	1	nil
String	24,16	nil	0	nil
Map	40,8	nil	0	nil
String	16	nil	4	nil
Float64	8	nil	4	nil

*simplified representation

Let's parse "[1,2]" into **&v** of type []int

stackIndex=0; tokenIndex=0

type ExpectType	offset,size uintptr	dest unsafe.Pointer	parent uint32
Slice	0,16	addr(v) 📌	4294967295
Int	0,8	nil	0

Set destination pointer and transition
to next stack frame

Let's parse "[1,2]" into **&v** of type []int

stackIndex=1; tokenIndex=1

type ExpectType	offset,size uintptr	dest unsafe.Pointer	parent uint32
Slice	0,16	addr(v)	4294967295
Int	8 🙌,8	addr(v)+offset(0) 🙌	0

Set destination pointer, write value, update offset and transition to next token

Let's parse "[1,2]" into **&v** of type []int

stackIndex=1; tokenIndex=2

type ExpectType	offset,size uintptr	dest unsafe.Pointer	parent uint32
Slice	0,16	addr(v)	4294967295
Int	16 🙌,8	addr(v)+offset(8) 🙌	0

Set destination pointer, write value, update offset and transition to next token

Let's parse "[1,2]" into **&v** of type []int

stackIndex=0; tokenIndex=3

type ExpectType	offset,size uintptr	dest unsafe.Pointer	parent uint32
Slice	0,16	addr(v)	4294967295
Int	8,8	addr(v)+offset(16)	0 🙌

Transition back to parent frame, last token
at root frame, return no error.

White House urges developers to dump C and C++

Biden administration calls for developers to embrace memory-safe programming languages and move away from those that cause buffer overflows and other memory access vulnerabilities.



By Grant Gross

InfoWorld | FEB 27, 2024 10:35 AM PST

Rust as an example of a programming language it considers safe. In addition, an NSA [cybersecurity information sheet](#) from November 2022 listed C#, [Go](#), Java, Ruby, and Swift, in addition to Rust, as programming languages it considers to be memory-safe.

```
type GoSlice struct {  
    Data      unsafe.Pointer  
    Length    int  
    Capacity int  
}
```

[1, 2]



Maps are complex!
No way around reflection
...right?

```
1 mapType := reflect.TypeOf(map[string]int)
2 vm := reflect.MakeMap(mapType)
3 key, value := "Meaning of life", 42
4 vm.SetMapIndex(reflect.ValueOf(key), reflect.ValueOf(value))
```

```
1
2 //go:linkname makemap reflect.makemap
3 func makemap(*typ, int) unsafe.Pointer
4
5 //nolint:golint
6 //go:linkname mapassign_faststr runtime.mapassign_faststr
7 //go:noescape
8 func mapassign_faststr(t *typ, m unsafe.Pointer, s string) unsafe.Pointer
9
10 //go:linkname mapassign runtime.mapassign
11 //go:noescape
12 func mapassign(t *typ, m unsafe.Pointer, k unsafe.Pointer) unsafe.Pointer
13
14 // typ represents reflect.rtype for noescape trick
15 type typ struct{}
16
17 type emptyInterface struct { _ *typ; ptr unsafe.Pointer }
18
19 func getTyp(t reflect.Type) *typ {
20     return (*typ)(((*emptyInterface)(unsafe.Pointer(&t))).ptr)
21 }
22
23 func canUseAssignFaststr(mapType reflect.Type) bool {
24     return mapType.Elem().Size() ≤ 128 && mapType.Key().Kind() = reflect.String
25 }
26
```



1	MakeMapInlined	49.8 ns/op	144 B/op	1 allocs/op
2	MakeMap	121.2 ns/op	400 B/op	3 allocs/op
3	MakeMapReflect	192.9 ns/op	560 B/op	5 allocs/op
4	MakeMapFastUnsafe	127.4 ns/op	400 B/op	3 allocs/op

```
1 // MakeMapInlined
2 m := make(map[string]int)
3 m["Meaning of life"] = 42
```

```

1  m := make(map[string]string, tokens[ti].Elements)
2  tiEnd := tokens[ti].End
3
4  for ti++; ti < tiEnd; ti += 2 {
5      tokVal := tokens[ti+1]
6      if tokVal.Type != jscan.TokenTypeString {
7          if tokVal.Type == jscan.TokenTypeNull {
8              key := s[tokens[ti].Index+1 : tokens[ti].End-1]
9              keyUnescaped := unescape.Valid[S, string](key)
10             m[keyUnescaped] = ""
11             continue
12         }
13         err = ErrorDecode{
14             Err: ErrUnexpectedValue,
15             Index: tokVal.Index,
16         }
17         return true
18     }
19     key := s[tokens[ti].Index+1 : tokens[ti].End-1]
20     keyUnescaped := unescape.Valid[S, string](key)
21     value := s[tokVal.Index+1 : tokVal.End-1]
22     m[keyUnescaped] = unescape.Valid[S, string](value)
23 }
24 ti++
25 (*map[string]string)(p) = m

```

Type `map[string]string`
 is very common,
 inline it as
`ExpectTypeMapStringString`

ROADMAP



Testing, testing, testing, benchmarking...



jscan v3 release (Tokenizer, Decoder, Unmarshal)



json/v2 options

Duplicates, InvalidUTF8, CaseInsensitiveFieldMatch



Further performance improvements



io.Reader support?

github.com/romshark/jscan

github.com/romshark/jscan-experimental-decoder